

Automation for the AI Factory

Reports: Foundry models and token

Model: GPT 5.1 - RG `gh-esml-project004-sdc-dev-001-rg` - window 30d - generated 2026-06-29 16:51

2) Current workload telemetry (from logs)

Workload - telemetry	Value	Unit
Total users with access	50000	users
INPUT tokens per minute	0	TPM
OUTPUT tokens per minute	0	TPM
REQUESTS per minute	0	RPM
Total tokens (window)	0	tokens
EA Discount	29%	
Cache rate	35%	
Input rate (\\$/1M)	0.27	
Cached rate (\\$/1M)	0.099	
Output rate (\\$/1M)	7.81	

3) Derived workload information

Derived	Value	Unit
Average TPM	0	TPM
Requests per day	0	req/day
Monthly tokens (30d)	0	tokens
Est. monthly PAYGO cost	0	USD

Recommendations

Based on PAYGO usage, is PTU an option

GPT 5.1 PTU Recommendation		Value	Note
PTUs (avg, ~35% cache)	0	4750 input TPM/PTU	
-> resulting TPM	0	normalized	
PTU to handle spikes	False	else spillover PAYGO	
AI Gateway loadbalancer	True		